


State U3 Unemployment Rate

The dataset, its clocks, and every validation check

Monthly state-level U3 (LAUS, seasonally adjusted) from FRED/ALFRED · source notebook: ur_validation_new.ipynb



Registered results: 10 PASS 5 WARN 0 FAIL

July 2026

One report, two datasets: how state U3 reaches the world

- **What it is** — `{ABBR}UR`: the state headline (U3) unemployment rate. Monthly, **seasonally adjusted**, percent, from the BLS **Local Area Unemployment Statistics (LAUS)** program, served by FRED with archival vintages in ALFRED. State labor force (`{ABBR}LF`) rides along for reconciliation; national **UNRATE** is the benchmark.
- **H1 · Publication schedule** — published in the BLS *State Employment and Unemployment* report, the **same monthly release as state payrolls**, on a **Friday ~3 weeks after** the national Employment Situation — about **50 days after the reference month**. Verified in V3.2–V3.4.
- **H2 · Reference period** — the rate reflects labor-force status in the **week containing the 12th** of the month; FRED stamps the observation on the **1st**. Verified in V2.1.
- **H3 · Revision structure** — LAUS is re-estimated and **benchmarked annually** each winter (re-anchored to CES/QCEW inputs, re-seasonally-adjusted). Revisions are moderate and **quantized to 0.1pp** — the rate is published to one decimal, so many true revisions round to exactly zero. Quantified in V4.

Note — U3 is the **third revision profile** of the state labor trio: claims are (almost) never revised, payrolls always, and U3 sits in between — **86% of first prints change**, but 84% of changes stay within ± 0.5 pp. The suite is tuned to a **bounded, quantized rate** rather than a level: range checks in $[0, 40]\%$, revisions in percentage points, and a digit-grid check instead of a magnitude guard.

Three tables, three clocks — same numbers, different questions

table	file	rows	key (grain)	span	question it answers
vintage (df_v)	ur_vintage_ first_release.csv	11,745	(state, obs_date)	obs 2005-06 → 2026-05 released 2005-07-22 → 2026-06-23	“What was first published for month t, and when? ”
current (ur_c)	ur_current_history.csv	30,804 = 51 × 604 months (the 'ur' slice; 'lf' stored alongside)	(state, obs_date)	obs 1976-01 → 2026-05	“What does FRED say today about month t?” (fully revised)
pit	ur_pit_features.csv	11,133 = ramp 7→51 states × 239 as-of months	(as_of_date, state)	as-of 2006-08 → 2026-06	“What did we actually know on day d? ” (backtest feed)

Note — The same physical number lives in all three tables — what differs is **when you are allowed to see it**. ``df_v`` runs on the **information clock** (what the world learned, and when), ``ur_c`` on the **observation clock** (the cleanest history, for targets and climatology), and ``pit`` replays both clocks so that **a backtest can only touch what was known at the time**. With 86% of U3 prints later revised, confusing them is not a rounding error.

Column dictionary — vintage & current tables

column	table	meaning
state / abbr / fips	both	full state name, USPS code, FIPS code (51 incl. DC)
obs_date	both	reference month, stamped on the 1st — the rate measures the survey week containing the 12th (V2.1)
series	current	'ur' = unemployment rate (validated) · 'lf' = labor force (auxiliary, same file)
value_current	current	today's fully-revised rate for that month (%)
first_release_date	vintage	day the value first entered the ALFRED archive — the information clock
value_first_release	vintage	the first print (%) — what a real-time observer saw
is_current	vintage	True if the first print still equals today's value (only ~14% of rows — V4.1)
pub_lag_days	vintage	first_release_date – obs_date: 45–111d, median 50d (V3.3)

Note — If you remember one thing: **obs_date says what the number measures; first_release_date says when anyone could know it.** For U3 the two clocks sit ~50 days apart — almost two months of information lag baked into every real-time use of this series.

Column dictionary — the point-in-time panel (20 columns)

column(s)	meaning
as_of_date	the vantage point : a month-start date; every cell in the row reflects only data released on or before it (239 consecutive month-starts)
state	one row per state per as-of date (ramps 7 → 51 as states enter the archive — see V3.1)
ur_t0 ... ur_t12	the 13 most recent monthly rates available as of that date; t0 = newest
ur_latest_obs	which observation month t0 refers to
ur_lag_months	as-of month – latest obs month: how stale t0 is (median 2, max 4 — V3.7)
ur_mom1_chg / ur_mom3_chg / ur_yoy_chg	1-month, 3-month and 12-month changes in pp , computed only from released readings

Discussion — “Vantage point” in plain English: as_of_date is the day you pretend it is *now*. Standing on the 1st of month M, the newest state U3 you legally have is month **M–2** (the ~50-day lag): the panel encodes that staleness in ur_lag_months instead of ever showing you a number early. Zero look-ahead, verified row-by-row in V3.7.

Why the three tables start on different dates

layer	starts	why it can't start earlier
current history	1976-01	FRED serves today's vintage across the whole history — no point-in-time obligation; LAUS reaches back to 1976 for every state
vintage (first prints)	2005-06 (first release 2005-07-22)	ALFRED began archiving state U3 in Jul-2005 — and only for a 7-state pilot (MO, MS, KY, IN, TN, IL, AR). The other 44 states join 2007-07-20 . A first print never snapshotted cannot be recovered (V3.1)
pit panel	2006-08	a row needs 13 monthly readings (t0...t12) → first valid as-of ≈ 14 months after a state's archive start; the full 51-state cross-section exists from mid-2008

- The ordering **current** ≥ **vintage** ≥ **pit** is structural — the same staircase as claims (1986 → 2010-09 → 2010-11) and NFP (1990 → 2005 → 2006).
- Consequence: **1976 → 2005 is NOT point-in-time**. Usable for climatology and cycle studies; never for backtesting (V3.1).

Note — Reading ``describe()`` output: the **50% row is the midpoint of the sample, not a boundary** — the pit panel's median as-of lands in 2017 only because the panel spans 2006→2026 and the 7-state early years contribute few rows. The **start** is the min row: 2006-08 (thin), full 51-state panel from 2008–09.

22 checks in 4 modules — how to read the rest of this deck

module	theme	checks	registered results
V1	Structural integrity — states, keys, nulls, bounded rates, 0.1pp grid, truncation	V1.0 – V1.5 (+ missing-value heatmaps)	print-style flags (OK)
V2	Calendar & frequency — the observation clock (the month being measured)	V2.1, V2.2a, V2.2b, V2.3	3 PASS · 1 WARN
V3	Point-in-time & release calendar — the information clock (when it became known)	V3.1 – V3.7	5 PASS · 2 WARN
V4	Revision behavior — first release and after (in percentage points)	V4.1 – V4.4	2 PASS · 2 WARN

- **PASS** = criterion met · **WARN** = usable with a documented caveat · **FAIL** = blocker. This run: **10 PASS, 5 WARN, 0 FAIL**.
- Every check states its **question and criterion up front**; each slide shows the evidence, the verdict, and — where it matters — what it means for modeling.
- Scope: this notebook is the per-question restructure of modules **V1–V4**; deeper modules (outliers, coherence, seasonality, lead-lag, backtest hygiene, governance: V5–V10) live in the full suite ``ur_validation.ipynb``.

Is the panel physically sound — right states, unique keys, complete fields, credible rates?

Structural integrity

check	question it answers
V1.0	What do the three tables actually look like? (eyeball before testing)
V1.1	Do we have all 51 states in every table, every period?
V1.2	Is (state, month) a unique key — no double-loaded rows?
V1.3	Are the key fields complete (no nulls)?
V1.4	Is every rate finite and inside a plausible bounded range (0–40%)?
V1.4b	Are the rates quantized to 0.1pp as published? Plus: what do the missing-value heatmaps really show?
V1.5	Did any state's fetch silently truncate?

V1.0 — The three tables at a glance

INFO

Criterion: look before testing — a validation suite that never prints its data validates blind.

- **vintage**: 11,745 rows · first prints **1.6% → 28.2%** — the max is Nevada's Apr-2020 COVID print, a useful sanity anchor · pub_lag 45–111d, median 50d.
- **current**: 30,804 rows = 51 × 604 months · obs **1976-01 → 2026-05** · revised rates span **1.7% → 30.5%**.
- **pit**: 11,133 rows · 239 month-start as-of dates · 20 columns of look-ahead-free features.

	rows	rate range	dates
vintage	11,745	1.6 .. 28.2%	2005-06 .. 2026-05
current	30,804	1.7 .. 30.5%	1976-01 .. 2026-05
pit	11,133	(t0..t12)	2006-08 .. 2026-06


Note — A rate is bounded and pre-normalized — no log transforms or per-capita scaling needed, unlike the two level datasets (claims, payrolls). The flip side: information arrives in **coarse 0.1pp steps** (V1.4b), which shapes everything V4 measures.

V1.1 — Do we have all 51 states in every table, every period?

OK

Criterion: 51 unique states per table; 51 per period outside documented edges.

- Unique states: **vintage 51 · current 51 · pit 51**. Per period: current **51 always** (1976→); vintage and pit dip to **7 states for 24 months (2005-06 → 2007-06)** — the archive pilot era, not missing data.




Note — The step from 7 → 51 in mid-2007 is ALFRED's archive coverage expanding, and it is the single explanation behind every “missing” artifact in this dataset (see the heatmap slide). The current table is flat at 51 because today's vintage covers all states over the full history.

V1.2 — Is (state, month) a unique key — no double-loaded rows?

OK

Criterion: 0 duplicate rows on the declared key of each table.

- **vintage 11,745 rows = 11,745 unique keys · current 30,804 = 30,804 · pit 11,133 = 11,133** (key: as_of_date × state) — zero duplicates anywhere.




Note — Duplicate keys are the classic symptom of a fetch retried and appended twice. Zero duplicates + V1.5's row-count check together rule out both double-loading and silent truncation.

V1.3 — Are the key fields complete (no nulls)?

OK

Criterion: 0 nulls in every check field of all three tables.

- Field-by-field: **0 nulls** in state, obs_date, first_release_date, value_first_release (vintage); value_current (current); as_of_date, state (pit).




Note — FRED encodes a missing month as the string '.' — the fetch layer converts and drops those before assembly, so a null here would mean a parser regression. Genuinely unpublished months appear as **absent rows** (V2.2's subject) — and U3 has exactly one, Oct-2025.

V1.4 — Is every rate finite and inside a plausible bounded range?

OK

Criterion: 0 rows outside [0, 40] %.

- **0 bad rows** in both tables. Observed range: vintage [1.6, 28.2]%, current [1.7, 30.5]% — the upper tail is Apr–Jun 2020 (NV, MI, HI), well inside the 40% guard.



Note — A bounded rate makes the plausibility check trivial compared with the level datasets — no digit-error guard needed: a decimal slip (4.5 → 45) would land outside the band and be caught here.


V1.4b — Are the rates quantized to 0.1pp, as published?

OK

Criterion: 100% of values are exact multiples of 0.1pp.

- 100.0000% of vintage and current values sit on the 0.1pp grid — 11 distinct decimal digits, near-uniformly used (no digit preference).

V1.4b — first-decimal digit distribution (confirms 0.1pp grid)





Note — This is why V4 looks the way it does: a true revision smaller than 0.05pp **rounds to exactly zero** in the published series. The 14% “never revised” share in V4.1 is partly a rounding illusion — the underlying re-estimation touches essentially every month.

V1.4b (cont.) — Missing-value heatmaps — one block, one explanation

INFO

Criterion: know what a black cell means before calling it missing data.

- vintage by **obs_date** → **1,056 “missing”** of 12,801 (8.25%) · by **first_release_date** → **1,060** (8.31%) · current → **0** · pit ur_t0 by as_of → **1,056** of 12,189 (8.66%).
- All four numbers are the same object: **44 states × 24 months (2005-06 → 2007-06) = 1,056** — the pre-archive block of the states that joined in Jul-2007.




Note — Different story from claims: there, the black cells came from **staggered capture** (post-2020 states archived on different days). Here capture is synchronized — one shared release day per month — so both pivots show the **same ramp-in block**: absence of archive coverage, not data loss inside covered range. The current table shows zero because today's vintage back-fills all history.

V1.5 — Did any state's fetch silently truncate?

OK

Criterion: max – min rows per state ≤ 3 months.

- Rows per state (current history): **604 = 604 = 604** — min, median and max identical, spread **0**. A perfectly rectangular 51 × 604 panel.



Note — Cleaner than claims (spread 12 from ragged 1986 starts): LAUS publishes every state from 1976-01, so any deviation from 604 would be a fetch failure, not history.

The observation clock: `obs_date` stamps the reference month — the rate measures the survey week containing the 12th.

Calendar & frequency


check	question it answers
V2.1	Do all observation dates follow the monthly (month-start) convention?
V2.2a	Are there missing months shared by all states (systematic BLS holes)?
V2.2b	Are there per-state (idiosyncratic) missing months?
V2.3	Do all states cover the same window (balanced panel)?

V2.1 — Do all observation dates follow the monthly month-start convention?

PASS

Criterion: 100% of obs_date have day == 1.

- 0 non-month-start rows in vintage and current alike.




Note — H2 anchors here: FRED stamps the month on the 1st, but the rate measures the **week containing the 12th** — pairing U3 with weekly claims must respect that reference week, not the stamp (the LAUS/claims alignment rule).

V2.2a — Are there missing months shared by all states (systematic BLS holes)?

WARN

Criterion: documented; shared BLS-wide holes tolerated as WARN.

- Exactly **one systematic hole: 2025-10**, absent for **all 51 states and for national UNRATE alike** — the late-2025 federal shutdown canceled the October household survey; the estimates were never produced.




Note — RULE: **do not interpolate across Oct-2025** — there is no true value to recover, and a filled value would leak the shape of the recovery into features. Pipelines must tolerate a NaN month; the pit builder simply goes one month staler through the gap (V3.7's max lag of 4).

V2.2b — Are there per-state (idiosyncratic) missing months?

PASS

Criterion: == 0; each occurrence must be investigated.

- **0 idiosyncratic missing months** — no state ever skipped a month on its own in 50 years of history.




Note — LAUS is a centrally-produced model estimate, not state self-reporting — so unlike claims (where a state can fail to transmit), a hole can only be BLS-wide. That is exactly the pattern observed: one shared hole, zero private ones.

V2.3 — Do all states cover the same window (balanced panel)?

PASS

Criterion: start spread == 0; end spread ≤ 1 month.

- Start spread **0** (every state begins 1976-01) · end spread **0** (every state ends 2026-05). Common balanced window: **1976-01 → 2026-05, 604 months**, no clipping needed.



Note — The joint monthly release keeps all 51 states in lockstep — there is no live capture edge like the weekly claims panel, because the whole report lands at once.

The information clock:
first_release_date stamps when the number became known. ~50 days after the month, with state detail 3 weeks behind the national report.

Point-in-time & release calendar


check	question it answers
V3.1	How far back can point-in-time be reconstructed?
V3.2	Does the national vintage calendar match the jobs-report (Friday) schedule?
V3.3	Is the state publication lag ~50 days?
V3.4	Do state releases follow the national jobs report?
V3.5	Is the publication-lag regime stable over time?
V3.6	Do months become available in order (monotone frontier)?
V3.7	Does the PIT panel contain any look-ahead?

V3.1 — How far back can point-in-time be reconstructed?

WARN

Criterion: boundary documented (WARN = limitation, not defect).

- First ALFRED vintage **2005-07-22** — but only for a 7-state pilot (MO, MS, KY, IN, TN, IL, AR); the remaining **44 states join 2007-07-20**. PIT panel reconstructable from 2006-08 (thin) / **full 51-state cross-section from mid-2008**.




Note — Everything left of the boundary — 29 years, 1976→2005 — exists only as **today's values**: usable for climatology and cycle studies, never for backtesting. For clean cross-sectional work, **2008-07 onward** is the honest start.

V3.2 — Does the national vintage calendar match the jobs-report (Friday) schedule?

PASS

Criterion: Friday ≥ 80% of national UNRATE vintage dates.

- 797 national UNRATE vintages: **Friday 84.2%**; the rest (Mon 2%, Tue 3%, Wed 4%, Thu 7%) are ALFRED late captures and occasional shifted releases.




Note — The national Employment Situation lands the **first Friday** after the reference month. State U3 then follows ~3 weeks later (V3.4) — a two-step information cascade a lead-lag study can exploit: the national rate is known long before its state decomposition.

V3.3 — Is the state publication lag ~50 days?

PASS

Criterion: median publication lag in [40, 60] days.

- 11,745 state first releases: **median 50d, IQR 48–53d**, range 45–111d — about **7 weeks** after the reference month, jointly with state payrolls.




Note — Compare the clocks across the trio: claims **13–19d**, U3/NFP **~50d**. By the time a state's U3 for month M prints, you already hold ~7 weekly claims readings covering the same period — the core of any claims→U3 nowcast.

V3.4 — Do state releases follow the national jobs report?

PASS

Criterion: $\geq 95\%$ of state releases after the national report for the same month.

- **100.0%** of state first releases land after the national jobs report; **median gap +14d** (IQR 14–18, max 70).




Note — State detail never front-runs the national aggregate — so the national print is legitimately in your information set when forecasting the state cross-section two weeks later. A useful conditioning variable, verified rather than assumed.

V3.5 — Is the publication-lag regime stable over time?

WARN

Criterion: yearly median lag range ≤ 7 days.

- 2005–2024: yearly medians pinned in **48.5–50.5d** — remarkably stable. Then **2025: 52d** and **2026: 66d** blow the range out to 18d — the shutdown-disrupted release calendar catching up.




Note — One rule for 19 years, then a live disruption: backtests through 2024 may assume a constant ~50d lag; **2025–26 need actual release dates**, not the rule. Expect the calendar to renormalize once the delayed reports catch up.

V3.6 — Do months become available in order (monotone frontier)?

PASS

Criterion: 0 release-order inversions per state.

- **0 inversions** — every state's release dates are weakly increasing in observation order.




Note — An inversion would mean a later month became visible before an earlier one — poison for “latest available” logic. Monotonicity is what lets the pit builder use a simple cummax **availability frontier** with no correction passes.

V3.7 — Does the PIT panel contain any look-ahead?

PASS

Criterion: 0 violations, 0 unmatched keys.

- Re-verified all **11,133 rows**: release date of the underlying value $\leq$ as_of_date — **0 violations, 0 unmatched keys**.
- Staleness (ur_lag_months): **median 2, min 2, max 4** — at any month-start you know U3 through roughly two months back.



Note — The hard gate for backtesting, recomputed **against the vintage table** rather than the builder's own bookkeeping — a builder bug could not certify itself. The stretch to 4 months of staleness is the Oct-2025 hole and the delayed 2025–26 calendar passing through the panel.

MODULE

V4

First release and after, measured in percentage points: how often, where in the calendar, with what bias — and how the national rate calibrates it.

Revision behavior


check	question it answers
V4.1	How often, and by how much, does a first print get revised (endpoint: first vs current)?
V4.2	Do the larger revisions cluster at the annual (winter) benchmark?
V4.3	Is there a systematic bias in first prints?
V4.4	How large is the national UNRATE revision (calibration)?

V4.1 — How often, and by how much, does a first print get revised?

WARN

Criterion: informational — U3 is moderately, quantized-revised.

- 11,745 obs compared: **86% revised** · 14% exact-zero · only 25% within $\pm 0.1\text{pp}$ · **84% within $\pm 0.5\text{pp}$** · max **[4.1pp]** (Wyoming Apr-2020: 9.2 $\rightarrow$ 5.1).
- The entire top-10 of largest revisions is **2020 COVID months** — first prints made in the fog, rewritten by the next benchmark.




Note — U3 is the middle of the spectrum: claims 3.33% revised, **U3 86%**, NFP 99%. And 86% **understates** the churn — with a 0.1pp grid, any re-estimate below 0.05pp rounds to zero (V1.4b). Practical reading: the typical revision is small ($\pm 0.5\text{pp}$ covers 84%), but **backtests still need vintage data** — benchmark years and crisis months move by whole points.

V4.2 — Do the larger revisions cluster at the annual (winter) benchmark?

PASS

Criterion: $\geq 50\%$ of largest revisions released Jan–Apr.

- **99%** of each obs-month's largest revision lands in a **Jan–Apr** release; modal month **March** (1,581 of 1,656 obs; CA/TX/NY full vintage trajectories).




Note — The annual LAUS benchmark — re-anchoring and re-seasonal-adjustment — lands each winter and rewrites history in one pass. Apply revisions **on their release date, never backward**: a backtest that uses the revised value before March of the following year is trading on the benchmark early.

V4.3 — Is there a systematic bias in first prints?

WARN

Criterion: $|mean\ revision| < 0.03pp$.

- Mean revision **-0.047pp** (median +0.00, std 0.45pp): first prints **overstate** the rate slightly on average, and the by-year chart shows the negative sign persisting through most of the 2010s.




Note — A small persistent sign is what a benchmark direction looks like — and a predictable error is a **modellable** one: a first-print-correction term (or simply expecting -0.05pp) is free accuracy for anyone consuming the vintage series in real time.

V4.4 — How large is the national UNRATE revision (calibration)?

PASS

Criterion: informational calibration on the national series.

- National UNRATE first-release → current: **median +0.00pp, IQR [+0.00, +0.10], max |0.4pp|** — an order of magnitude tamer than the state max of 4.1pp.



Note — The national rate is revised on the same annual schedule but from a far larger sample — it anchors expectations: state revisions ~10× the national tail. If your model consumes both, give the national rate the higher real-time trust.

Scorecard — every registered check

ID	check	result	headline evidence
V2.1	All obs_date are month-starts	PASS	0 non-month-start rows
V2.2a	Systematic (shared) missing months	WARN	one BLS-wide hole: Oct-2025 (shutdown; national UNRATE too)
V2.2b	Idiosyncratic (per-state) missing months	PASS	none in 50 years
V2.3	Panel edges aligned across states	PASS	start spread 0 · end spread 0 (1976-01 → 2026-05)
V3.1	ALFRED vintage boundary identified	WARN	2005-07 (7-state pilot) / 2007-07 (all 51); pre-2005 climatology-only
V3.2	National vintages land on jobs-report Friday	PASS	Friday 84.2%
V3.3	State publication lag ~50 days	PASS	median 50d (IQR 48–53)
V3.4	State releases follow the national report	PASS	100% after national; median gap +14d
V3.5	Publication-lag regime stable	WARN	48–50d for 19 years; 2025–26 disrupted (52d, 66d)
V3.6	Release dates weakly increasing per state	PASS	0 inversions
V3.7	PIT panel look-ahead guard	PASS	0 violations / 0 unmatched of 11,133 rows; staleness 2–4m
V4.1	Endpoint revision rate & magnitude	WARN	86% revised; 84% within ±0.5pp; max 4.1pp (WY 2020-04)
V4.2	Revisions cluster at the annual benchmark	PASS	99% of largest revisions in Jan–Apr; modal March
V4.3	Systematic bias in first prints	WARN	mean −0.047pp — first prints slightly overstate
V4.4	National UNRATE revision calibration	PASS	median +0.00pp, max 0.4pp — 10× tamer than states

Module V1 uses print-style flags rather than registered rows: V1.1 OK · V1.2 OK · V1.3 OK · V1.4 OK (bounded, COVID max 28.2%) · V1.4b OK (100% on the 0.1pp grid) · V1.5 OK (perfect 51×604 rectangle).

What a model built on this data must respect

- **Backtests must use vintage data.** 86% of first prints change; the current history embeds annual benchmarks that were unknown in real time. The pit panel (0 look-ahead, staleness 2–4 months) is the only table a backtest may touch.
- **The information clock is slow.** Median 50-day lag → standing on any month-start you know U3 through M–2. Weekly claims cover the gap — that is the nowcasting opportunity, not a defect.
- **The point-in-time world starts 2005-07 for 7 pilot states, 2007-07 for all 51** — treat mid-2008 as the first clean cross-section. 1976→2005 is climatology only.
- **Benchmark season is Jan–Apr, modal March.** 99% of the biggest revisions land there; apply them on release date, never backward. First prints carry a small persistent overstatement (–0.047pp) — a modellable correction.
- **Respect the 0.1pp grid.** Sub-0.05pp changes are invisible; treat the rate as discrete (25% of revisions round to within ± 0.1 pp) and don't over-interpret single-tick moves.
- **Never interpolate across Oct-2025.** The BLS-wide hole has no true value; the 2025–26 release calendar is disrupted (median lag 52d → 66d) — use actual release dates there.

Note — Verdict: **10 PASS · 5 WARN · 0 FAIL** — cleared for forecasting, lead-lag and backtesting under the rules above. The WARNs document the data's plumbing (archive ramp-in, shutdown hole, benchmark bias), not defects; deeper modules (outliers, coherence, seasonality, lead-lag orientation, backtest hygiene, governance) run in the full suite `ur_validation.ipynb` as V5–V10.